

DPA Assignment 3

Girish Rajani

2023-03-05

#Recitation Problems

#Chapter 6

#1.

- (a) The best subset selection will have the lowest training RSS because it fits all P_k models that contain exactly k predictors. All combinations of p are taken into consideration here. This is not the case with forward and backward stepwise selection because they are greedy and don't consider all predictors.
- (b) For smallest test RSS, it is difficult to say based on the given information because best subset selection considers all predictors and has lowest training RSS so it may also have lowest testing RSS, however, it may overfit. The other two models, since they avoid overfitting, may also find a model that gives the lowest testing RSS.

#(c)

i - True - In forward stepwise selection, we add one additional parameter in M_k and so we select the best model within the predictors of the $(k+1)$ variable

ii - True - In backward stepwise selection, we use the subset of the predictors in the $(k+1)$ variable model and remove one predictor

iii - False - The predictors in backward stepwise selection are only identified by a subset of predictors in the $k+1$ variable model by the backward stepwise selection because the model will contain all but one predictor, therefore both models end up selecting different predictors

iv - False - The predictors in forward stepwise selection are only identified by a subset of predictors in the $k+1$ variable model by the forward stepwise selection because the model will contain one additional predictor, therefore both models end up selecting different predictors

v - False - The best subset selection predictors are not identified by a subset from the $k+1$ variable. It fits all k variables and contains exactly k predictors

#2.

(a)

- iii. Since Lasso yields sparse models meaning only a subset of variables is used so the model is less flexible. Additionally, as our λ increases and we

perform shrinkage and variable selection, the variance significantly decreases more than the increase in bias so we have a bias/variance trade off which gives improved prediction accuracy.

(b)

- iii. Ridge regression follows the same concept as above but we only perform shrinkage so our features shrink close to 0 but not exactly to 0. This means that as p goes closer to n , the increase in bias is less than the decrease in variance.

(c)

- ii. Non-linear models have high flexibility which means that the variance increases and so for improved predictions, we will have a significant decrease in bias and a slight increase in variance.

#3. (a) iv. The training RSS will steadily decrease as we increase s from 0 because as s increases, the least square solution falls within this budget and will become less restrictive. Therefore, the training RSS will steadily decrease until it yields this solution.

(b)

- ii. The testing RSS will initially decrease but as it fits the model, flexibility increases which will cause overfitting to occur and result in an increase in RSS.

(c)

- iii. The variance will steadily increase as our model flexibility increases

(d)

- iv. The bias will steadily decrease as our model flexibility increases

(e)

- v. The irreducible error is not dependent on s therefore it remains constant

#4. (a) iii. The training RSS will steadily increase because an increase in λ will also increase the penalty term

(b)

- ii. The testing RSS will initially decrease but as it fits the model, flexibility decreases which will cause overfitting to occur and result in an increase in RSS.

(c)

- iv. The variance will steadily decrease as our model flexibility decreases

(d)

- iii. The bias will steadily increase as our model flexibility decreases

(e)

- v. The irreducible error is not dependent on λ therefore it remains constant

5) a

$$\text{minimize}_{\beta} \left\{ \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2 \right\}$$

When $\hat{\beta}_0 = 0$, $n=2$, $p=2$, in this setting, we get:

$$\text{minimize } \left\{ (y_1 - \hat{\beta}_1 x_{11} - \hat{\beta}_2 x_{12})^2 + (y_2 - \hat{\beta}_1 x_{21} - \hat{\beta}_2 x_{22})^2 + \lambda (\hat{\beta}_1^2 + \hat{\beta}_2^2) \right\}$$

b We can assume that since $x_1 \leq x_2$, we can sub them as x_1
 x_2 x_3 x_4 x_5 x_6 x_7

Rewrite $a \rightarrow$

$$\text{minimize} \left((y_1 - \hat{\beta}_1 x_1 - \hat{\beta}_2 x_2)^2 + (y_2 - \hat{\beta}_1 x_2 - \hat{\beta}_2 x_2)^2 + \lambda(\hat{\beta}_1^2 + \hat{\beta}_2^2) \right)$$

where we can take partial derivative w.r.t β_1 and equate to 0:

$$\hat{\beta}_1(x_1^2 + x_2^2 + \lambda) + \hat{\beta}_2(x_1^2 + x_2^2) - y_1 x_1 - y_2 x_2 = 0$$

We can take partial derivative w.r.t β_2 and equte to 0:

$$\hat{\beta}_1(x_1^2 + x_2^2) + \hat{\beta}_2(x_1^2 + x_2^2 + \lambda) - y_1x_1 - y_2x_2 = 0$$

$$\hat{\beta}_1(x_1^2 + x_2^2 + \lambda) + \hat{\beta}_2(x_1^2 + x_2^2) = y_1 x_1 + y_2 x_2 = \hat{\beta}_1(x_1^2 + x_2^2) + \hat{\beta}_2(x_1^2 + x_2^2 + \lambda) - y_1 x_1 - y_2 x_2$$

Simplify:

$$\hat{\beta}_1 + \hat{\beta}_2(x_1^2 + x_2^2) = \hat{\beta}_2(x_1^2 + x_2^2) + \hat{\beta}_2$$

$$\hat{\beta}_1 X = \hat{\beta}_2 X$$

$$\beta_1 = \hat{\beta}_2$$

c) Lasso regression Optimization:

$$\text{minimize } \left\{ \sum_{i=1}^n \left(y_i - \hat{\beta}_0 - \sum_{j=1}^p \hat{\beta}_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\hat{\beta}_j| \right\}$$

When $\hat{\beta}_0 = 0$, $n=2$, $p=2$, in this setting we get:

$$\text{minimize } \left\{ (y_1 - \hat{\beta}_1 x_{11} - \hat{\beta}_2 x_{12})^2 + (y_2 - \hat{\beta}_1 x_{21} - \hat{\beta}_2 x_{22})^2 + \lambda (|\hat{\beta}_1| + |\hat{\beta}_2|) \right\}$$

d) Replace the penalty term from Lasso regression and derivate w.r.t β

$$\frac{\partial}{\partial \beta} (\lambda |\beta|) = \frac{\lambda |\beta|}{\beta}$$

We also expand this for β_1 & β_2 :

$$\frac{\lambda |\beta_1|}{\beta_1} = \frac{\lambda |\beta_2|}{\beta_2}$$

Assuming that both β_1 and β_2 are either positive or negative.

5c and 5d

#Chapter 7


#2. (a) As lambda goes to infinity, the penalty term is significant, $g(x) = 0$ and hence $\hat{g} = 0$

(b) As lambda goes to infinity, and $\hat{g}^1(x) = 0$, then \hat{g} will be a constant

- (c) As lambda goes to infinity and $m = 2$, $g^2(x) = 0$ which will result in a straight line, a function $ax+b$
- (d) As lambda goes to infinity and $m = 3$, $g^3(x) = 0$ which will result in a quadratic plot, a function $ax^2 + bx + c$
- (e) As lambda is 0, there will be no penalty term applied to \hat{g} . This means that the straight line will be drawn through all points perfectly

#3.

```
library(ggplot2)
x <- -2:2
y <- 1+x-2*(x-1)^2*I(x>1)
plot(x, y, type="o")
```



From above, we can see that when X goes from -2 to 1, the estimated curve is linear which results in a slope and y intercept of 1. However, when X is greater than one, we can observe a quadratic estimated curve.


#4.

```
x <- -2:2
y <- c(1 + 0 + 0, # x = -2
      1 + 0 + 0, # x = -1
```

```

1 + 1 + 0, # x = 0
1 + (1-0) + 0, # x = 1
1 + (1-1) + 0 # x = 2
)
plot(x,y, type = "o")

```



The estimated curve from above shows a slope of -2

#5. (a) As lambda goes to infinity, \hat{g}_2 will have a smaller training RSS because it is a more flexible model. It is more flexible because it has a higher penalty term of the fourth derivative order when compared to \hat{g}_1 which has a lower penalty term.

(b) As lambda goes to infinity, \hat{g}_1 is more likely to have a smaller training RSS because it is less likely to overfit due to its lower derivative penalty term. Since \hat{g}_2 is more flexible and has lower training RSS, it may overfit on test RSS.

(c) Since lambda is 0 here, there will be no penalty term applied in neither models. This means that both models will have the same train/test RSS.

```

#Import necessary libraries
data("mtcars")
data("swiss")
library(glmnet)

## Loading required package: Matrix

## Loaded glmnet 4.1-6

```

```

library("readxl")
library(mgcv)

## Loading required package: nlme

## This is mgcv 1.8-41. For overview type 'help("mgcv-package")'.

library(visreg)

#Problem 1

#Load mtcars sample dataset into dataframe
mtcars <- data.frame(mtcars)
set.seed(5)

#Perform a basic 80/20 test-train split
split <- sample(c(rep(0, 0.8 * nrow(mtcars)), rep(1, 0.2 * nrow(mtcars))))
train <- mtcars[split == 0, ]
test <- mtcars[split == 1, ]

summary(mtcars)

##           mpg           cyl           disp           hp
##  Min.      :10.40   Min.      :4.000   Min.      : 71.1   Min.      : 52.0
##  1st Qu.:15.43   1st Qu.:4.000   1st Qu.:120.8   1st Qu.: 96.5
##  Median :19.20   Median :6.000   Median :196.3   Median :123.0
##  Mean     :20.09   Mean     :6.188   Mean     :230.7   Mean     :146.7
##  3rd Qu.:22.80   3rd Qu.:8.000   3rd Qu.:326.0   3rd Qu.:180.0
##  Max.     :33.90   Max.     :8.000   Max.     :472.0   Max.     :335.0
##           drat           wt           qsec           vs
##  Min.      :2.760   Min.      :1.513   Min.      :14.50   Min.      :0.0000
##  1st Qu.:3.080   1st Qu.:2.581   1st Qu.:16.89   1st Qu.:0.0000
##  Median :3.695   Median :3.325   Median :17.71   Median :0.0000
##  Mean     :3.597   Mean     :3.217   Mean     :17.85   Mean     :0.4375
##  3rd Qu.:3.920   3rd Qu.:3.610   3rd Qu.:18.90   3rd Qu.:1.0000
##  Max.     :4.930   Max.     :5.424   Max.     :22.90   Max.     :1.0000
##           am           gear           carb
##  Min.      :0.0000   Min.      :3.000   Min.      :1.000
##  1st Qu.:0.0000   1st Qu.:3.000   1st Qu.:2.000
##  Median :0.0000   Median :4.000   Median :2.000
##  Mean     :0.4062   Mean     :3.688   Mean     :2.812
##  3rd Qu.:1.0000   3rd Qu.:4.000   3rd Qu.:4.000
##  Max.     :1.0000   Max.     :5.000   Max.     :8.000

head(mtcars)

##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160 110 3.90 2.620 16.46  0  1    4    4
## Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02  0  1    4    4
## Datsun 710     22.8   4  108  93 3.85 2.320 18.61  1  1    4    1
## Hornet 4 Drive 21.4   6  258 110 3.08 3.215 19.44  1  0    3    1

```



```
## Hornet Sportabout 18.7    8   360 175 3.15 3.440 17.02  0  0    3    2
## Valiant            18.1    6   225 105 2.76 3.460 20.22  1  0    3    1

cat("Dimensions of training data: ", dim(train))

## Dimensions of training data:  26 11

cat("\nDimensions of testing data: ", dim(test))

##
## Dimensions of testing data:  6 11

#fit a linear model using mpg as target response
lm <- lm(mpg ~ ., data=train)
summary (lm)

##
## Call:
## lm(formula = mpg ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.0238 -1.6448 -0.2359  0.7270  5.6116
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 17.87836   20.11527   0.889   0.388
## cyl         -0.40777    1.14174  -0.357   0.726
## disp         0.01432    0.02184   0.656   0.522
## hp          -0.02233    0.02514  -0.888   0.388
## drat        -0.78630    2.43324  -0.323   0.751
## wt          -3.56516    2.23639  -1.594   0.132
## qsec         0.72054    0.80297   0.897   0.384
## vs           0.51670    2.27364   0.227   0.823
## am           1.87919    2.35810   0.797   0.438
## gear         1.44036    1.85758   0.775   0.450
## carb        -0.16122    0.96871  -0.166   0.870
##
## Residual standard error: 2.698 on 15 degrees of freedom
## Multiple R-squared:  0.8643, Adjusted R-squared:  0.7739
## F-statistic: 9.555 on 10 and 15 DF,  p-value: 7.622e-05
```

#Explain what features are relevant based on t-statistic #A smaller t-value indicates high similarity/relationship between two features.

#Based on the t-value observations from the linear model fit, we can say that features hp and wt (with wt being the most relevant) are selected as the most relevant features.

```
#What are the associated coefficient values for relevant features?
lm$coefficients
```



```
## (Intercept)          cyl          disp          hp          drat          wt
## 17.87835971 -0.40776819  0.01432014 -0.02233375 -0.78629698 -3.56515667
##          qsec          vs          am          gear          carb
##  0.72054326  0.51670020  1.87919113  1.44035680 -0.16121614
```

#The out-of-sample test set performance (using predict) on the regular linear model:

```
y_hat <- predict(lm, test, type = "response")
y <- test[, "mpg"]
MSE <- mean((y_hat - y)^2)
cat("OOS MSE on testing data: ", MSE)
```

```
## OOS MSE on testing data:  8.823345
```

#Perform a ridge regression using glmnet

#pass x matrix and y vector:

```
x <- model.matrix(mpg ~ ., data=train)[, -1]
y <- train$mpg
```

#Perform ridge regression using 100 values of Lambda for tuning

```
grid <- 10^seq(10, -2, length = 100)
ridge.mod <- cv.glmnet(x, y, alpha = 0, lambda = grid, parallel=TRUE)
```


```
## Warning: executing %dopar% sequentially: no parallel backend registered
```

```
## Warning: Option grouped=FALSE enforced in cv.glmnet, since < 3
observations per
## fold
```

```
min_lambda <- ridge.mod$lambda.min
cat("Minimum value of Lambda: ", min_lambda)
```

```
## Minimum value of Lambda:  3.511192
```

#Plot training MSE as a function of Log Lambda
plot(ridge.mod)



```
# Fitting Ridge Regression with optimal Lambda
ridge.mod2 <- glmnet(x, y, alpha = 0, lambda = min_lambda)
```

```
coef(ridge.mod2)
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              s0
## (Intercept) 21.58560850
## cyl        -0.41252847
## disp       -0.00530685
## hp         -0.01037774
## drat        0.75754158
## wt         -1.08253237
## qsec        0.13080551
## vs          0.86473981
## am          1.24208873
## gear        0.63172887
## carb       -0.45619319
```

#By observing the coefficients in the regular linear model and the ridge regression model from above, the coefficients have come closer to zero. Since none of the coefficients are zero, ridge regression here does not perform variable selection but instead performs shrinkage.

```
#pass x matrix on testing
```

```
x_test <- model.matrix(mpg~. ,test)[-1]
```

```
# Predict out-of-sample test set performance
y_hat_test <- predict(ridge.mod2, s = min_lambda, newx = x_test)
y_test <- test[, "mpg"]
test_MSE <- mean((y_hat_test - y_test)^2)
```

```
cat("Out-of-Sample test set performance MSE: ", test_MSE)
```

```
## Out-of-Sample test set performance MSE: 10.14883
```

#By observing the out of sample test set performance MSE from the regular linear model (3.97) and the ridge regression model (2.33), the MSE has decreased which shows that the ridge regression model has performed better. After observing the coefficients and the MSE of both models, we can say that the ridge regression model has performed shrinkage.

#Problem 2

```
#Load swiss sample dataset into dataframe
```

```
swiss <- data.frame(swiss)
set.seed(5)
```

```
#Perform a 80/20 test-train split
```

```
split <- sample(c(rep(0, 0.8 * nrow(swiss)), rep(1, 0.2 * nrow(swiss))))
train <- swiss[split == 0, ]
test <- swiss[split == 1, ]
```

```
summary(swiss)
```

```
##      Fertility      Agriculture      Examination      Education
##  Min.   :35.00    Min.   : 1.20    Min.   : 3.00    Min.   : 1.00
## 1st Qu.:64.70    1st Qu.:35.90    1st Qu.:12.00   1st Qu.: 6.00
##  Median :70.40    Median :54.10    Median :16.00   Median : 8.00
##  Mean   :70.14    Mean   :50.66    Mean   :16.49   Mean   :10.98
## 3rd Qu.:78.45    3rd Qu.:67.65    3rd Qu.:22.00   3rd Qu.:12.00
##  Max.   :92.50    Max.   :89.70    Max.   :37.00   Max.   :53.00
##      Catholic      Infant.Mortality
##  Min.   : 2.150    Min.   :10.80
## 1st Qu.: 5.195    1st Qu.:18.15
##  Median :15.140    Median :20.00
##  Mean   :41.144    Mean   :19.94
## 3rd Qu.:93.125    3rd Qu.:21.70
##  Max.   :100.000    Max.   :26.60
```

```
head(swiss)
```

```
##      Fertility Agriculture Examination Education Catholic
## Courtelary      80.2      17.0      15      12      9.96
## Delemont       83.1      45.1       6       9      84.84
## Franches-Mnt   92.5      39.7       5       5      93.40
## Moutier        85.8      36.5      12       7      33.77
## Neuveville     76.9      43.5      17      15       5.16
```

```
## Porrentruy      76.1      35.3      9      7      90.57
## Infant.Mortality
## Courtelary      22.2
## Delemont        22.2
## Franches-Mnt    20.2
## Moutier         20.3
## Neuveville      20.6
## Porrentruy      26.6
```

```
cat("Dimensions of training data: ", dim(train))
```

```
## Dimensions of training data: 38 6
```

```
cat("\nDimensions of testing data: ", dim(test))
```

```
##
```

```
## Dimensions of testing data: 9 6
```

```
#fit a linear model using Fertility as target response
```

```
lm <- lm(Fertility ~ ., data=train)
```

```
summary(lm)
```

```
##
```

```
## Call:
```

```
## lm(formula = Fertility ~ ., data = train)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -15.8999  -4.8833   0.3608   4.4918  14.6768
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   66.13831    11.75847   5.625 3.23e-06 ***
## Agriculture   -0.15836     0.07813  -2.027 0.051092 .
## Examination   -0.26886     0.32164  -0.836 0.409404
## Education     -0.82147     0.21968  -3.739 0.000724 ***
## Catholic       0.10987     0.04499   2.442 0.020312 *
## Infant.Mortality 1.05061     0.41608   2.525 0.016722 *
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 7.297 on 32 degrees of freedom
```

```
## Multiple R-squared:  0.7224, Adjusted R-squared:  0.679
```

```
## F-statistic: 16.65 on 5 and 32 DF, p-value: 4.262e-08
```

#Explain what features are relevant based on t-statistic #A smaller t-value indicates similarity/relationship between two features.

#Based on the t-value observations from the linear model fit, we can say that features Education and Agriculture (with Education being the most relevant) are selected as the most relevant features.

#What are the associated coefficient values for relevant features?

```
lm$coefficients
```

```
##      (Intercept)      Agriculture      Examination      Education
##      66.1383136      -0.1583572      -0.2688619      -0.8214705
##      Catholic Infant.Mortality
##      0.1098728      1.0506114
```

#The out-of-sample test set performance (using predict) on the regular linear model:

```
y_hat <- predict(lm, test, type = "response")
```

```
y <- test[, "Fertility"]
```

```
MSE <- mean((y_hat - y)^2)
```

```
cat("OOS MSE on testing data: ", MSE)
```

```
## OOS MSE on testing data: 45.75983
```

#Perform a Lasso regression using glmnet

#pass x matrix and y vector:

```
x <- model.matrix(Fertility ~ ., data=train)[, -1]
```

```
y <- train$Fertility
```

#Perform Lasso regression using 100 values of lambda for tuning

```
grid <- 10^seq(10, -2, length = 100)
```

```
lasso.mod <- cv.glmnet(x, y, alpha = 1, lambda = grid, parallel=TRUE)
```


```
min_lambda <- lasso.mod$lambda.min
```

```
cat("Minimum value of Lambda: ", min_lambda)
```

```
## Minimum value of Lambda: 0.09326033
```

#Plot training MSE as a function of Log Lambda

```
plot(lasso.mod)
```



```
# Fitting Lasso Regression with optimal Lambda
lasso.mod2 <- glmnet(x, y, alpha = 1, lambda = min_lambda)

coef(lasso.mod2)

## 6 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept)  65.0640375
## Agriculture  -0.1417047
## Examination  -0.2608186
## Education     -0.7940812
## Catholic       0.1057848
## Infant.Mortality 1.0498219
```

#By observing the coefficients in the regular linear model and the lasso regression model from above, there isn't much difference, however, some shrinkage is observed.

#Typically, with lasso, we would expect to see shrinkage and then variable selection occurring. However, since none of the coefficients have shrunk to 0, we can say that our lasso model only performs shrinkage and not both shrinkage and variable selection.

```
#pass x matrix on testing
x_test <- model.matrix(Fertility~. ,test)[,-1]

# Predict out-of-sample test set performance
y_hat_test <- predict(lasso.mod2, s = min_lambda, newx = x_test)
y_test <- test[, "Fertility"]
```

```
test_MSE <- mean((y_hat_test - y_test)^2)

cat("Out-of-Sample test set performance MSE: ", test_MSE)

## Out-of-Sample test set performance MSE: 46.35392
```

#By observing the out of sample test set performance MSE from the regular linear model (49.85) and the lasso regression model (47.95), the MSE has slightly decreased which shows that the lasso regression model has performed slightly better.

#Problem 3

#Load Concrete Compressive Strength sample dataset from UCI Repository into dataframe using readxl package

```
set.seed(5)
CCS_data <- read_excel("D:\\Users\\giris\\Downloads\\Concrete_Data.xls")
colnames(CCS_data) <- c("C1", "C2", "C3", "C4", "C5", "C6", "C7", "C8", "CCS")
```

```
summary(CCS_data)
```

```
##           C1           C2           C3           C4
##  Min.      :102.0   Min.       : 0.0   Min.       : 0.00   Min.       :121.8
## 1st Qu.:192.4   1st Qu.:  0.0   1st Qu.:  0.00   1st Qu.:164.9
## Median :272.9   Median : 22.0   Median :  0.00   Median :185.0
## Mean   :281.2   Mean   : 73.9   Mean   : 54.19   Mean    :181.6
## 3rd Qu.:350.0   3rd Qu.:142.9   3rd Qu.:118.27   3rd Qu.:192.0
## Max.    :540.0   Max.    :359.4   Max.    :200.10   Max.    :247.0
##           C5           C6           C7           C8
##  Min.       : 0.000   Min.       :801.0   Min.       :594.0   Min.       : 1.00
## 1st Qu.: 0.000   1st Qu.: 932.0   1st Qu.:731.0   1st Qu.:  7.00
## Median : 6.350   Median : 968.0   Median :779.5   Median : 28.00
## Mean   : 6.203   Mean   : 972.9   Mean   :773.6   Mean   : 45.66
## 3rd Qu.:10.160   3rd Qu.:1029.4   3rd Qu.:824.0   3rd Qu.: 56.00
## Max.    :32.200   Max.    :1145.0   Max.    :992.6   Max.    :365.00
##           CCS
##  Min.       : 2.332
## 1st Qu.:23.707
## Median :34.443
## Mean   :35.818
## 3rd Qu.:46.136
## Max.    :82.599
```

```
head(CCS_data)
```

```
## # A tibble: 6 × 9
##       C1      C2      C3      C4      C5      C6      C7      C8      CCS
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  540      0      0    162    2.5 1040    676     28   80.0
## 2  540      0      0    162    2.5 1055    676     28   61.9
## 3  332.   142.      0    228     0    932    594    270   40.3
## 4  332.   142.      0    228     0    932    594    365   41.1
```



```
## 5 199. 132.    0 192  0   978. 826.  360 44.3
## 6 266  114    0 228  0   932  670   90 47.0
```

#Using gam function to predict the CCS as a non-linear function of the input features C1-C6

```
GAM_non_linear <- gam(CCS ~ C1+C2+C3+C4+C5+C6, data=CCS_data)
summary(GAM_non_linear)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## CCS ~ C1 + C2 + C3 + C4 + C5 + C6
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5.326997  10.510518   0.507 0.612387
## C1           0.108256   0.005214  20.761 < 2e-16 ***
## C2           0.079357   0.006193  12.814 < 2e-16 ***
## C3           0.055928   0.009287   6.022 2.4e-09 ***
## C4          -0.103871   0.027796  -3.737 0.000197 ***
## C5           0.356016   0.110251   3.229 0.001281 **
## C6           0.008027   0.006272   1.280 0.200940
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## R-sq.(adj) =  0.445   Deviance explained = 44.9%
## GCV = 155.83   Scale est. = 154.77    n = 1030
```

#With an adjusted R^2 of 0.445 from this model, only 45% of the variance is explained in this model which means that there is not a strong correlation between the predictors and response variables.

#Creating a GAM with linear terms as well as smoothed terms

```
GAM_linear <- gam(CCS ~ s(C1)+s(C2)+s(C3)+s(C4)+s(C5)+s(C6), data=CCS_data)
summary(GAM_linear)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## CCS ~ s(C1) + s(C2) + s(C3) + s(C4) + s(C5) + s(C6)
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  35.8178    0.3566  100.4 <2e-16 ***
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##          edf Ref.df      F  p-value
## s(C1) 4.464  5.513 69.530 < 2e-16 ***
## s(C2) 2.088  2.578 48.091 < 2e-16 ***
## s(C3) 5.332  6.404  1.784  0.101
## s(C4) 8.567  8.936 13.504 < 2e-16 ***
## s(C5) 7.133  8.143  5.498 1.22e-06 ***
## s(C6) 1.000  1.000  0.018  0.892
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.531   Deviance explained = 54.4%
## GCV = 134.84   Scale est. = 130.96      n = 1030
```

#With an adjusted R^2 of 0.531, this model suggests that a little bit more than half of the variance in the outcome data is explained by the model

```
#Compute R^2 value for initial non linear model
GAM_non_linear.sse <- sum(fitted(GAM_non_linear) - CCS_data$CCS)^2
GAM_non_linear.ssr <- sum(fitted(GAM_non_linear) - mean(CCS_data$CCS))^2
GAM_non_linear.sst <- GAM_non_linear.sse + GAM_non_linear.ssr

GAM_non_linear.r2 <- 1 - (GAM_non_linear.sse/GAM_non_linear.sst)

cat("R-square value of initial GAM non linear model:", GAM_non_linear.r2)

## R-square value of initial GAM non linear model: 0.4967177


#Compute R^2 value for the linear model with smoothing
GAM_linear.sse <- sum(fitted(GAM_linear) - CCS_data$CCS)^2
GAM_linear.ssr <- sum(fitted(GAM_linear) - mean(CCS_data$CCS))^2
GAM_linear.sst <- GAM_linear.sse + GAM_linear.ssr

GAM_linear.r2 <- 1 - (GAM_linear.sse/GAM_linear.sst)


cat("R-square value of GAM with smoothing:", GAM_linear.r2)

## R-square value of GAM with smoothing: 0.5000744


#Visualize the regression
visreg(GAM_non_linear, 'C1')
```




```
visreg(GAM_linear, 'C1')
```




```
visreg(GAM_non_linear, 'C2')
```




```
visreg(GAM_linear, 'C2')
```




```
visreg(GAM_non_linear, 'C3')
```




```
visreg(GAM_linear, 'C3')
```




```
visreg(GAM_non_linear, 'C4')
```




```
visreg(GAM_linear, 'C4')
```




```
visreg(GAM_non_linear, 'C5')
```




```
visreg(GAM_linear, 'C5')
```



```
visreg(GAM_non_linear, 'C6')
```

```
visreg(GAM_linear, 'C6')
```



#The confidence interval in the above visualizations is represented by grey band around the expected value (blue line).

#After having observed all vizualizations, it can be seen that after applying smoothing, the confidence interval is better and has higher values when compared to the initial model.